

Million Songs-Output#1

January 23, 2021

```
[3]: # %% [markdown]
# # Million Songs Problem Statement
#
# ### Context
# With the advent of technology, societies have become more efficient with
→their lives. But at the same time, individual human lives have become much
→more fast paced and distracted by leaving little time to explore artistic
→pursuits. Also, the technology has made significant advancements in the
→ability to coexist with art and general entertainment. In fact, it has made
→it easier for humans with shortage of time to find and consume good content.
→Therefore, one of the key challenges for the companies is to be able to
→figure out what kind of content their customers are most likely to consume.
→Almost every internet based company's revenue relies on the time consumers
→spend on their platforms. These companies need to be able to figure out what
→kind of content is needed in order to increase the time spent by customers
→on their platform and make their experience better.
# Spotify is one such audio content provider who has got a huge market base
→across the world. It has grown significantly because of its ability to
→recommend the 'best' next song to each and every customer based on the huge
→preference database they have gathered over time like millions of customers
→and billions of songs. This is done by using smart recommendation systems
→that can recommend songs based on the users' likes/dislikes
#
# ### Problem Statement
#
# Build a recommendation system to propose the top 10 songs for a user based on
→the likelihood of listening to those songs.
#
# ### Data Dictionary
# The core data is the Taste Profile Subset released by The Echo Nest as part
→of the Million Song Dataset. There are two files in this dataset. One
→contains the details about the song id, titles, release, artist name and the
→year of release. Second file contains the user id, song id and the play
→count of users.
#
# #### song_data
# 1. song_id - A unique id given to every song
```

```

# 2. title - Title of the song
# 3. Release - Name of the released album
# 4. Artist_name - Name of the artist
# 5. year - Year of release
#
# #### count_data
# 1. user_id - A unique id given to the user
# 2. song_id - A unique id given to the song
# 3. play_count - Number of times the song was played
#
# #### Data Source
# http://millionsongdataset.com/
#

```

```

[4]: from Functions_Million_Songs import *
from Functions_Million_Songs import tmp_pivot_table, sparse_matrix, dense_matrix
from Functions_Million_Songs import interactions_matrix      # made from ↵
      ↪ dfu_small

%load_ext autoreload
%autoreload 2
#%reload_ext autoreload

```

```

[5]: init_datsets()
from Functions_Million_Songs import dfs, dfu, dfus, dfu_small  #song, user and ↵
      ↪ joint datasets
dfs.info(), dfu.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 999056 entries, 0 to 999999
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   song_id         999056 non-null  object
1   title           999041 non-null  object
2   release         999051 non-null  object
3   artist_name     999056 non-null  object
4   year            999056 non-null  int64
dtypes: int64(1), object(4)
memory usage: 45.7+ MB
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2000000 entries, 0 to 1999999
Data columns (total 3 columns):
#   Column      Dtype
---  -
0   user_id     object

```

```
1  song_id      object
2  play_count   int64
dtypes: int64(1), object(2)
memory usage: 61.0+ MB
```

```
[5]: (None, None)
```

```
[6]: print(f'\n\n Only a few songs have been listened to a lot by a specific user,
      ↳so we keep these "outliers": \n')
dfu[['user_id', 'play_count']].nlargest(10, 'play_count'), dfs.count()

print_stats()

#Some Song titles appear in more than one song
print('\n\n Duplicate Titles example: "Intro"')
dfs.loc[dfs.title=='Intro']

print('\n\nSample data for "Kings Of Leon": \n ')
dfus.loc[dfus.artist_name=='Kings Of Leon'].sample(n=5)

top_n = 10
show_top_lists(top_n)

# surprise library -- Data Distributions and
# model evaluation (need more time to finish)
X = surprise_distribution_model_evaluation(dfus, dfu_small)
print(X)
# print(results)
```

Only a few songs have been listened to a lot by a specific user, so we keep these "outliers":

```
Some Listner and Song Stats:
    Total Users: 76353
    Total Artists: 72652
    Total Songs: 999056
    Total Song Titles: 702350
Total Albums Released: 149211
    Oldest Song Year: 1922
    Newest Song Year: 2011
    User-Song Density: 0.0026%
```

```
Duplicate Titles example: "Intro"
```

Sample data for "Kings Of Leon":

Top 10 Most Active Listeners

	user_id	play_count
Rank		
1	6d625c6557df84b60d90426c0116138b617b9449	711
2	fbee1c8ce1a346fa07d2ef648cec81117438b91f	643
3	4e11f45d732f4861772b2906f81a7d384552ad12	556
4	24b98f8ab023f6e7a1c37c7729c623f7b821eb95	540
5	1aa4fd215aadb160965110ed8a829745cde319eb	533
6	b04e41133dd3d30a5631cc8589a1eadd48a8bd53	523
7	15eeb36ae1c62d60de9fdeea0d121eb7d08713be	522
8	a15075a926c1998d91940f118342ba8356efc7d4	502
9	ce5c912bb8044f23fc0fc31bd986b8d0a7303db5	489
10	6a9cf03dfb2fc82f5b3b043c9c3fdbab997fd54d	487

Top 10 Most Listened Artist:

	artist_name	play_count
Rank		
1	Coldplay	70138
2	Kings Of Leon	68570
3	Florence + The Machine	60066
4	Dwight Yoakam	54136
5	Björk	53814
6	The Black Keys	52220
7	Jack Johnson	44083
8	Justin Bieber	41645
9	OneRepublic	40981
10	Train	39279

Top 10 most Played Songs:

	title \
Rank	
1	You're The One
2	Undo
3	Revelry
4	Horn Concerto No. 4 in E flat K495: II. Romanc...
5	Sehr kosmisch
6	Dog Days Are Over (Radio Edit)
7	Secrets
8	Canada
9	Invalid

10 Ain't Misbehavin

	artist_name	play_count
Rank		
1	Dwight Yoakam	54136
2	Björk	49253
3	Kings Of Leon	41418
4	Barry Tuckwell/Academy of St Martin-in-the-Fie...	31153
5	Harmonia	31036
6	Florence + The Machine	26663
7	OneRepublic	22100
8	Five Iron Frenzy	21019
9	Tub Ring	19645
10	Sam Cooke	18309

Top 10 Songs with most number of listeners:

	title \
Rank	
1	Sehr kosmisch
2	Undo
3	Dog Days Are Over (Radio Edit)
4	You're The One
5	Revelry
6	Secrets
7	Horn Concerto No. 4 in E flat K495: II. Romanc...
8	Fireflies
9	Hey_ Soul Sister
10	Tive Sim

	artist_name	Number_of_Listeners \
Rank		
1	Harmonia	8277
2	Björk	7032
3	Florence + The Machine	6949
4	Dwight Yoakam	6412
5	Kings Of Leon	6145
6	OneRepublic	5841
7	Barry Tuckwell/Academy of St Martin-in-the-Fie...	5385
8	Charttraxx Karaoke	4795
9	Train	4758
10	Cartola	4548

	song_id
Rank	
1	SOFRQTD12A81C233C0

2	SOAUWYT12A81C206F1
3	SOAXGDH12A8C13F8A1
4	SOBONKR12A58A7A7E0
5	SOSXLTC12AF72A7F54
6	SONYKOW12AB01849C9
7	SOEGIYH12A6D4FC0E3
8	SOLFXTT12AB017E3E0
9	SODJWHY12A8C142CCE
10	SOFLJQZ12A6D4FADA6

	song_id	count
2220	SOFRQTD12A81C233C0	8277
317	SOAUWYT12A81C206F1	7032
352	SOAXGDH12A8C13F8A1	6949
614	SOBONKR12A58A7A7E0	6412
7416	SOSXLTC12AF72A7F54	6145
...
8747	SOWNLZF12A58A79811	51
4492	SOLIGVL12AB017DBAE	51
622	SOBPGWB12A6D4F7EF3	50
9638	SOYYBJJ12AB017E9FD	48
2666	SOGSPGJ12A8C134FAA	48

[10000 rows x 2 columns]

```

Estimating biases using als...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using als...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using als...
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.
Computing the msd similarity matrix...
Done computing similarity matrix.

```

Computing the msd similarity matrix...
 Done computing similarity matrix.
 Computing the msd similarity matrix...
 Done computing similarity matrix.
 Computing the msd similarity matrix...
 Done computing similarity matrix.
 Computing the msd similarity matrix...
 Done computing similarity matrix.
 Estimating biases using als...
 Estimating biases using als...
 Estimating biases using als...

	test_rmse	fit_time	test_time
Algorithm			
BaselineOnly	5.384732	1.958337	3.614825
KNNBaseline	5.387468	15.956658	70.437691
KNNWithZScore	5.398714	12.596456	71.014284
CoClustering	5.417637	12.846680	3.501397
SlopeOne	5.420629	9.018290	23.741897
KNNWithMeans	5.428411	12.115448	65.781161
NMF	5.557056	42.861590	3.565619
KNNBasic	5.597561	11.745206	63.213809
NormalPredictor	6.641988	1.059326	3.476840
SVD	8.439450	37.268810	2.902564
SVDpp	8.442772	587.291683	27.076039

MODEL #1 - Collaborative Filtering - Using Cosine Similarity of User_ids ### dfus dataset(all users & songs data combo) ### It is possible to use altenate route(not implmented here) of ### just using top songs and users

```
[8]: init_model_1_CF_matrices()
from Functions_Million_Songs import tmp_pivot_table, sparse_matrix, dense_matrix

# pick user_id and # of song recommendations desired
user_id_to_recommend_songs_to = '01845f57f5c8b3309233e5a4a7145a7d33ad3d52'
num_of_songs_to_recommend = 5

# Show songs user has already listened to
x = pd.DataFrame(dfus.query('user_id ==_
↳@user_id_to_recommend_songs_to')[['title','artist_name']])
print(f'\n\nSongs {user_id_to_recommend_songs_to} has already listened to:\n_
↳{x[["title","artist_name"]]}')

# Recommend songs:
# find ordinal position of user_id index in tmp_pivot_table(user_id,song_id,_
↳play_count)
user_id_index = tmp_pivot_table.index.get_loc(
    tmp_pivot_table[tmp_pivot_table.index ==_
↳user_id_to_recommend_songs_to]
```

```

        .iloc[-1].name)
# get cosine similarity based recommendations
recommended_song_id_column_index = (recommendations_cosine_model_1
                                   (user_id_index,
                                   num_of_songs_to_recommend,
                                   dense_matrix))

print(f'\n\nTop {num_of_songs_to_recommend} song recommendations for_
→ "{user_id_to_recommend_songs_to}": \n')

get_song_ids_from_index(recommended_song_id_column_index, tmp_pivot_table)

```

Songs 01845f57f5c8b3309233e5a4a7145a7d33ad3d52 has already listened to:

	title \
1702697	Orgelblut
1702698	Welk
1702699	In These Arms
1702700	Die Kunst der Fuge_ BWV 1080 (2007 Digital Rem...
1702701	Don't Start Me Talkin'
1702702	Creil City
1702703	Together Again (Jimmy Jam Deep Remix)
1702704	Silent Shout
1702705	Marshall Examines His Carcass
1702706	Streets On Lock
1702707	Le Jardin d'Hiver
1702708	Schwarze Biene (Black Maja)
1702709	Représente
1702710	Love Is Not A Fight
1702711	The Gift
1702712	Trash
1702713	Kill Me
1702714	What We Do (Explicit) (Feat. Memphis Bleek)


	artist_name
1702697	Bohren & Der Club Of Gore
1702698	Bohren & Der Club Of Gore
1702699	Bon Jovi
1702700	Lionel Rogg
1702701	Sonny Boy Williamson
1702702	Alliance Ethnik
1702703	Janet Jackson
1702704	The Knife
1702705	Octopus Project
1702706	Young Jeezy
1702707	Jacky Terrasson

1702708	Bohren & Der Club Of Gore
1702709	Alliance Ethnik
1702710	Warren Barfield
1702711	Angels and Airwaves
1702712	The New York Dolls
1702713	Make the Girl Dance
1702714	Sauce Money Featuring Memphis Bleek

Top 5 song recommendations for "01845f57f5c8b3309233e5a4a7145a7d33ad3d52":

```
[8]: ['Make Love To Your Mind',
      'En Algún Lugar Del Puerto (2001 Digital Remaster)',
      'Sinisten tähtien alla',
      'Forever & Always',
      'How Long']
```

```
[9]: ## MODEL #2 - SVD Matrix Factorization
      ##### Split data into train,validate and test sets on dfus dataset(all users &
      ↳songs data combo)
      ##### same user_id may have been randomly split across train,validate and test
      ↳sets
      ##### It is possible to use alternate route(not implemented here) of keeping
      ##### user_id & song_id of a user either in train,validate and test but not
      ↳split across
      ##### these three sets
      X_train, X_test = train_test_split(dfus_small, test_size=0.2, random_state=42)
      X_train, X_val = train_test_split(X_train, test_size=0.1, random_state=42)
      u_train, s_train, vt_train, u_val, vt_val, u_test, vt_test =
      ↳perform_svd_test(X_train, X_val, X_test)
```



0.1 Accuracy Results:

For Train set the accuracy dips and then seems to increase as new features are added. The intersection of train &

The Graph is behaving oddly. For Validate & Test sets the accuracy is going down with adding more latent features

test sets suggests 450 might be optimal # of latent features to use.